



About Dataset

Context

The data contains information from the 1990 California census. It does provide an accessible introductory dataset for teaching people about the basics of machine learning.

Content

The data pertains to the houses found in a given California district and some summary stats about them based on the 1990 census data. The columns are as follows:

1. longitude: A measure of how far west a house is; a higher value is farther west
2. latitude: A measure of how far north a house is; a higher value is farther north
3. housingMedianAge: Median age of a house within a block; a lower number is a newer building
4. totalRooms: Total number of rooms within a block
5. totalBedrooms: Total number of bedrooms within a block
6. population: Total number of people residing within a block
7. households: Total number of households, a group of people residing within a home unit, for a block
8. medianIncome: Median income for households within a block of houses (measured in tens of thousands of US Dollars)
9. medianHouseValue: Median house value for households within a block (measured in US Dollars)
10. oceanProximity: Location of the house w.r.t ocean/sea

In [134]:

```
# Load needed libraries
import numpy as np
```

```

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import fetch_california_housing

california_df = fetch_california_housing(as_frame=True)
california_df = pd.read_csv("housing[1].csv")

print("Data Imported Successfully")

```

Data Imported Successfully

In [135...]: # Get the first 5 rows of data
california_df.head(5)

Out[135...]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population
0	-122.23	37.88	41.0	880.0	129.0	
1	-122.22	37.86	21.0	7099.0	1106.0	2
2	-122.24	37.85	52.0	1467.0	190.0	
3	-122.25	37.85	52.0	1274.0	235.0	
4	-122.25	37.85	52.0	1627.0	280.0	

In [136...]: # Get the last 5 rows of data
california_df.tail(5)

Out[136...]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population
20635	-121.09	39.48	25.0	1665.0	374.0	
20636	-121.21	39.49	18.0	697.0	150.0	
20637	-121.22	39.43	17.0	2254.0	485.0	
20638	-121.32	39.43	18.0	1860.0	409.0	
20639	-121.24	39.37	16.0	2785.0	616.0	

In [137...]: # Some info about the dataset
california_df.info()

||ANALYSIS||
We see that the total_bedrooms has a few NULL(MISSING) values. And the variab

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   longitude        20640 non-null   float64
 1   latitude         20640 non-null   float64
 2   housing_median_age 20640 non-null   float64
 3   total_rooms      20640 non-null   float64
 4   total_bedrooms   20433 non-null   float64
 5   population       20640 non-null   float64
 6   households       20640 non-null   float64
 7   median_income    20640 non-null   float64
 8   median_house_value 20640 non-null   float64
 9   ocean_proximity  20640 non-null   object  
dtypes: float64(9), object(1)
memory usage: 1.6+ MB

```

In [138...]: # Some description about the dataset
`california_df.describe()`

```

# ||ANALYSIS||
# Just a few columns dicating the mean, standard deviation and showing the 25th
# As well the smallest and largest value in each category

```

Out[138...]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms
count	20640.000000	20640.000000	20640.000000	20640.000000	20433.
mean	-119.569704	35.631861	28.639486	2635.763081	537.
std	2.003532	2.135952	12.585558	2181.615252	421.
min	-124.350000	32.540000	1.000000	2.000000	1.
25%	-121.800000	33.930000	18.000000	1447.750000	296.
50%	-118.490000	34.260000	29.000000	2127.000000	435.
75%	-118.010000	37.710000	37.000000	3148.000000	647.
max	-114.310000	41.950000	52.000000	39320.000000	6445.

In [139...]: # Get the Null values in our dataset
`pd.DataFrame(california_df.isnull().sum())`

```

# ||ANALYSIS||
# total_bedrooms has 207 missing values, we shall deal with this below

```

Out[139...]	0
	longitude 0
	latitude 0
	housing_median_age 0
	total_rooms 0
	total_bedrooms 207
	population 0
	households 0
	median_income 0
	median_house_value 0
	ocean_proximity 0

```
In [140...]: # Fix the total_bedrooms, replace the NULL values with the MEAN or AVERAGE val
total_bedrooms_mean_value = california_df["total_bedrooms"].mean()
california_df["total_bedrooms"].replace(np.nan, total_bedrooms_mean_value, inplace=True)
print("Successfully changed all NULL values to the mean value")
```

Successfully changed all NULL values to the mean value

/tmp/ipython-input-3536752185.py:3: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
california_df["total_bedrooms"].replace(np.nan, total_bedrooms_mean_value, inplace = True)
```

```
In [141...]: # Calculating mean of every column
column_count = 0
for i in california_df.columns:
    print(f"The mean value of {i} = {california_df[i].mean()}")
# Incrementing the Counter
column_count += 1
if column_count == 9: # I have used 9 because there are 10 columns
    break
```

```
The mean value of longitude = -119.56970445736432
The mean value of latitude = 35.63186143410853
The mean value of housing_median_age = 28.639486434108527
The mean value of total_rooms = 2635.7630813953488
The mean value of total_bedrooms = 537.8705525375617
The mean value of population = 1425.4767441860465
The mean value of households = 499.5396802325581
The mean value of median_income = 3.8706710029069766
The mean value of median_house_value = 206855.81690891474
```

Visualization

I have analysed and visualised many relationships and connections between different columns to get interesting insights from our dataset.

What have I done?

I have Explored the correlation & relationship between different data columns.

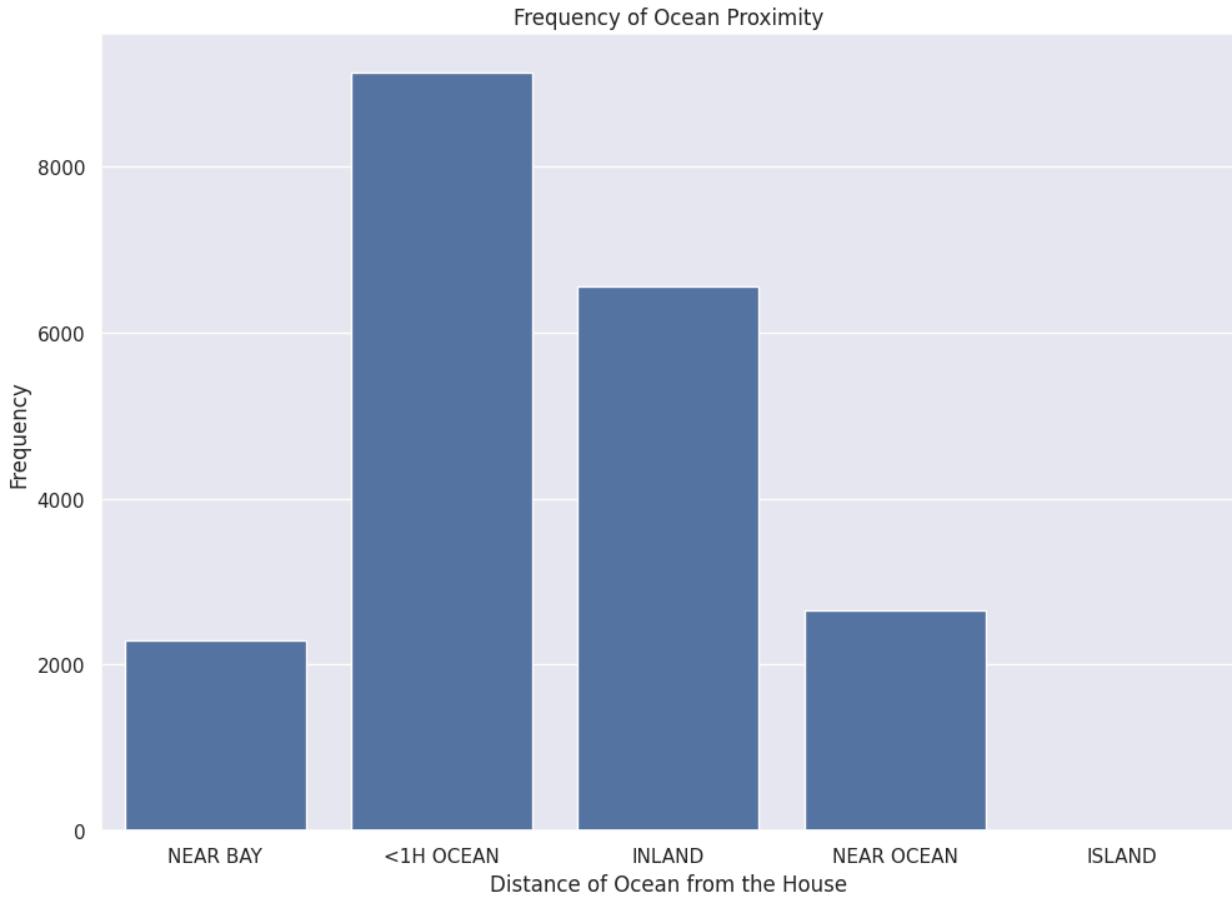
I have noted a few of the interesting insights from the exploratory analysis.

UNIVARIATE ANALYSIS

```
In [142]: # Frequency of Ocean Proximity
sns.set(rc={'figure.figsize':(11.7,8.27)})

# Plot
sns.countplot(data=california_df, x="ocean_proximity").set(
    title = "Frequency of Ocean Proximity",
    xlabel = "Distance of Ocean from the House",
    ylabel = "Frequency"
)

# Show
plt.show()
```



```
In [143]: for col in california_df.columns:  
    print("\n\n=====\nCOLUMN:", col)  
    print("=====\n")  
  
    # Check type  
    if pd.api.types.is_numeric_dtype(california_df[col]):  
        print("Type: Numeric")  
        print("Missing values:", california_df[col].isnull().sum())  
        print("Mean:", california_df[col].mean())  
        print("Median:", california_df[col].median())  
        print("Std:", california_df[col].std())  
        print("Skew:", california_df[col].skew())
```

```
=====
COLUMN: longitude
=====
Type: Numeric
Missing values: 0
Mean: -119.56970445736432
Median: -118.49
Std: 2.003531723502581
Skew: -0.2978012079524362

=====
COLUMN: latitude
=====
Type: Numeric
Missing values: 0
Mean: 35.63186143410853
Median: 34.26
Std: 2.1359523974571117
Skew: 0.46595300370997006

=====
COLUMN: housing_median_age
=====
Type: Numeric
Missing values: 0
Mean: 28.639486434108527
Median: 29.0
Std: 12.585557612111637
Skew: 0.060330637599136865

=====
COLUMN: total_rooms
=====
Type: Numeric
Missing values: 0
Mean: 2635.7630813953488
Median: 2127.0
Std: 2181.615251582787
Skew: 4.147343450632158

=====
COLUMN: total_bedrooms
=====
Type: Numeric
Missing values: 0
Mean: 537.8705525375617
Median: 438.0
Std: 419.26659232552385
Skew: 3.4770233756335105
```

```
=====
COLUMN: population
=====
Type: Numeric
Missing values: 0
Mean: 1425.4767441860465
Median: 1166.0
Std: 1132.4621217653375
Skew: 4.93585822672712

=====
COLUMN: households
=====
Type: Numeric
Missing values: 0
Mean: 499.5396802325581
Median: 409.0
Std: 382.3297528316099
Skew: 3.410437711667147

=====
COLUMN: median_income
=====
Type: Numeric
Missing values: 0
Mean: 3.8706710029069766
Median: 3.5347999999999997
Std: 1.8998217179452732
Skew: 1.6466567021344465

=====
COLUMN: median_house_value
=====
Type: Numeric
Missing values: 0
Mean: 206855.81690891474
Median: 179700.0
Std: 115395.6158744132
Skew: 0.9777632739098341

=====
COLUMN: ocean_proximity
=====
```

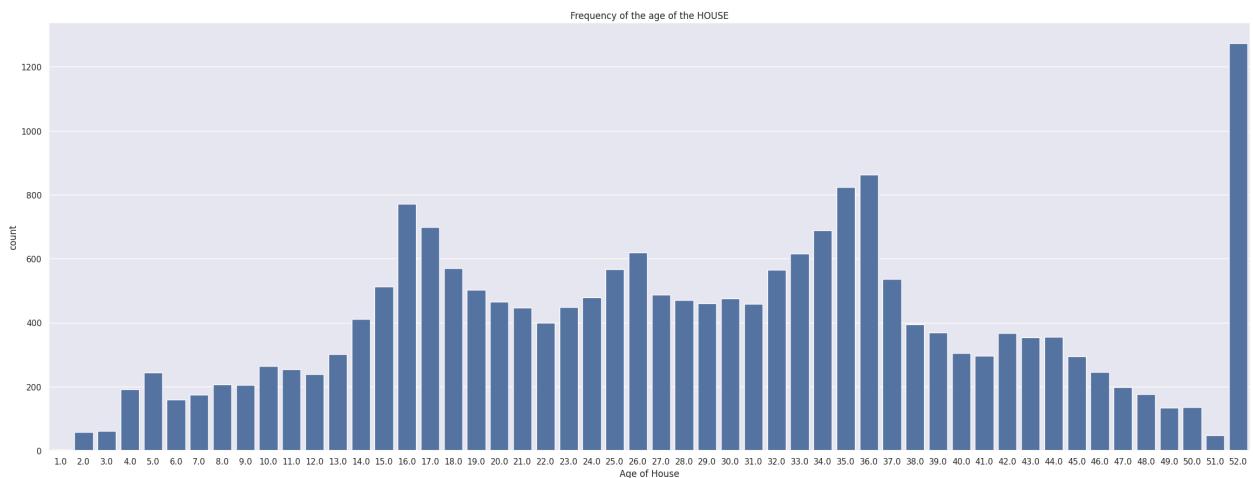
```
In [144...]: # Frequency of the age of the HOUSE
sns.set(rc={'figure.figsize':(28.7,10.27)})  
# Plot
```

```

sns.countplot(data=california_df, x="housing_median_age").set(
    title = "Frequency of the age of the HOUSE",
    xlabel = "Age of House"
)

# Show
plt.show()

```



```

In [145...]: print("===== CATEGORICAL SUMMARY =====\n")

categorical_cols = california_df.select_dtypes(exclude=[np.number]).columns.to

for col in categorical_cols:
    print(f"\n### {col} ###\n")
    print(california_df[col].value_counts(dropna=False))

```

===== CATEGORICAL SUMMARY =====

```

### ocean_proximity ###

ocean_proximity
<1H OCEAN      9136
INLAND         6551
NEAR OCEAN     2658
NEAR BAY        2290
ISLAND          5
Name: count, dtype: int64

```

```

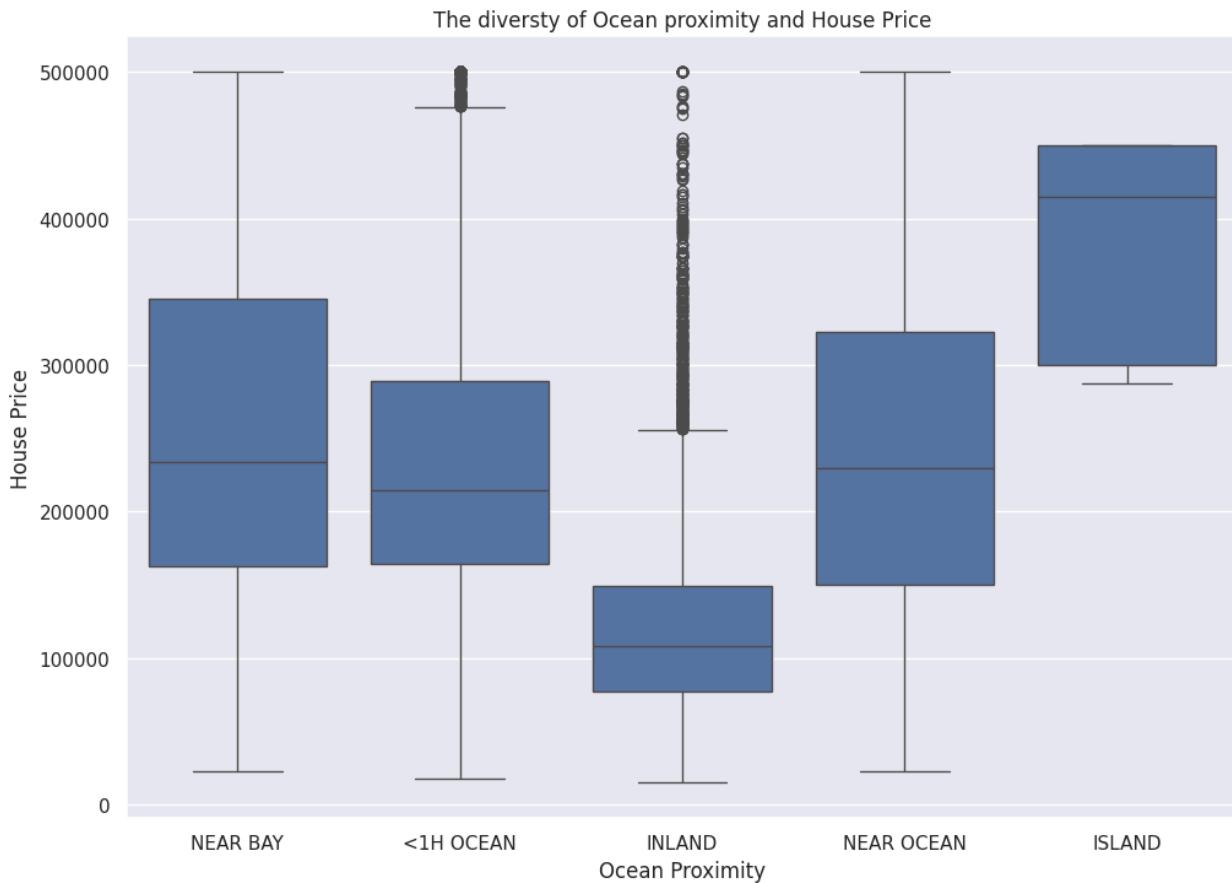
In [146...]: # The diversity of Ocean proximity and House Price
sns.set(rc={'figure.figsize':(11.7,8.27)})

sns.boxplot(x='ocean_proximity',y='median_house_value',data=california_df).set(
    title = "The diversity of Ocean proximity and House Price",
    xlabel = "Ocean Proximity",
    ylabel = "House Price"
)

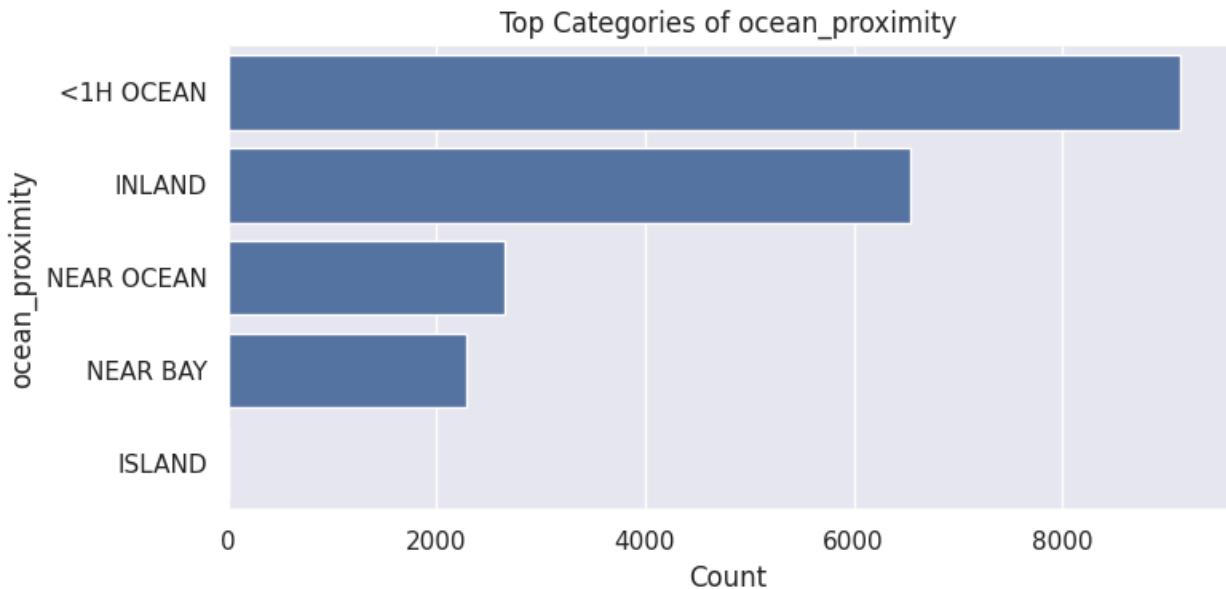
# Show

```

```
plt.show()
```



```
In [147]: for col in categorical_cols:  
    print(f"\n\n### {col} ###")  
    vc = california_df[col].value_counts().head(20)  
  
    plt.figure(figsize=(8,4))  
    sns.barplot(x=vc.values, y=vc.index)  
    plt.title(f"Top Categories of {col}")  
    plt.xlabel("Count")  
    plt.ylabel(col)  
    plt.tight_layout()  
    plt.show()  
  
### ocean_proximity ###
```



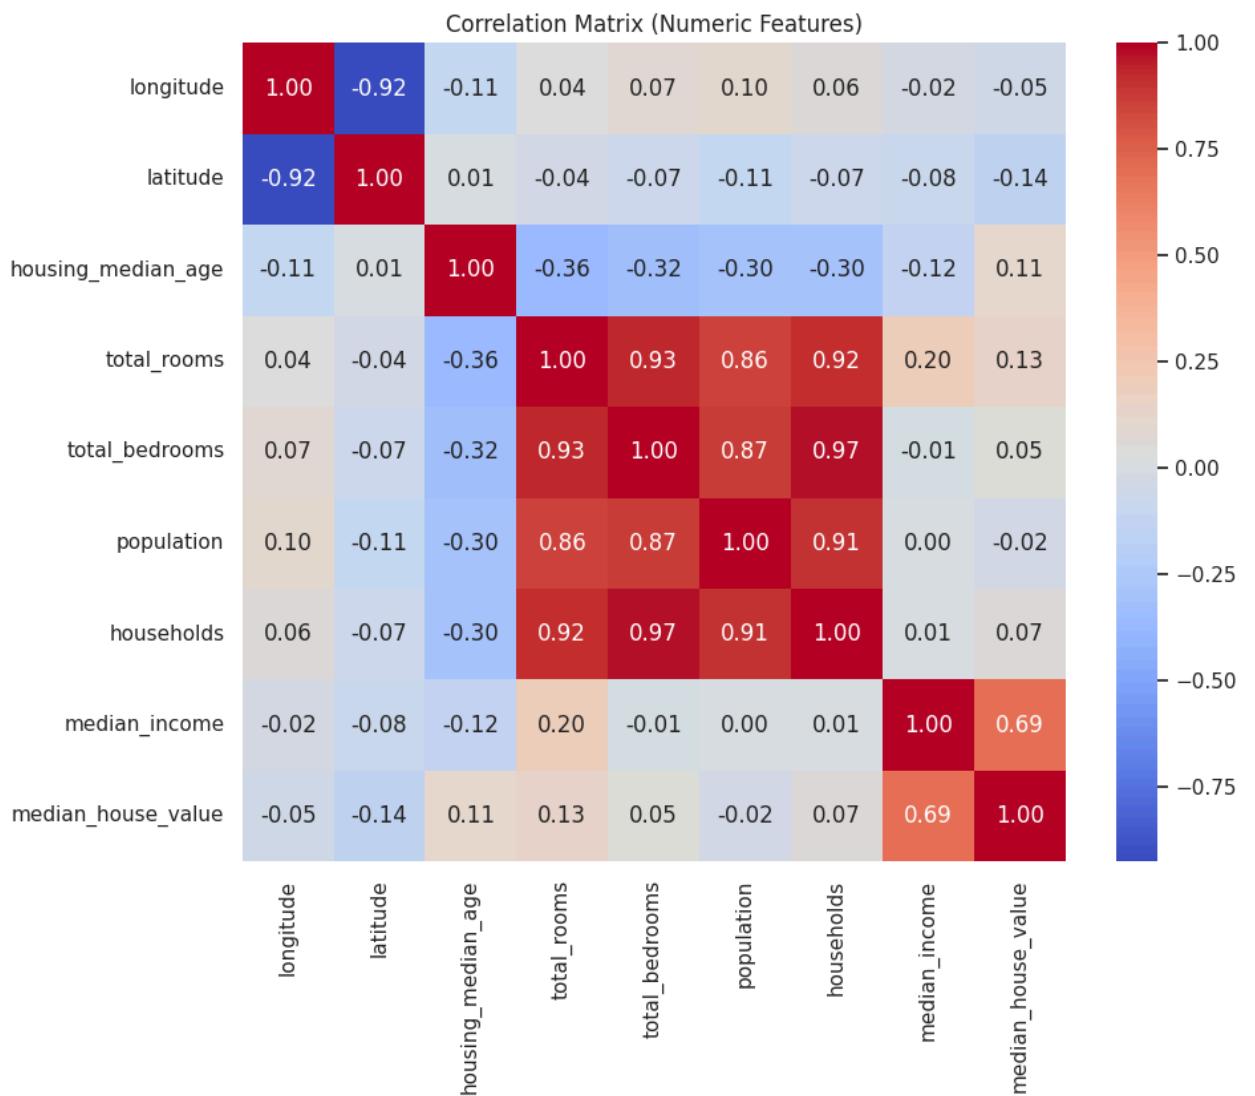
BI VARIATE ANALYSIS

```
In [148]: numeric_cols = california_df.select_dtypes(include=[np.number]).columns.tolist()
categorical_cols = california_df.select_dtypes(exclude=[np.number]).columns.tolist()

print("Numeric columns:", numeric_cols)
print("Categorical columns:", categorical_cols)

Numeric columns: ['longitude', 'latitude', 'housing_median_age', 'total_rooms',
'total_bedrooms', 'population', 'households', 'median_income', 'median_house_value']
Categorical columns: ['ocean_proximity']
```

```
In [149]: plt.figure(figsize=(10,8))
sns.heatmap(california_df[numeric_cols].corr(), annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Correlation Matrix (Numeric Features)")
plt.show()
```



```
In [150]: corr = california_df[numeric_cols].corr().abs()

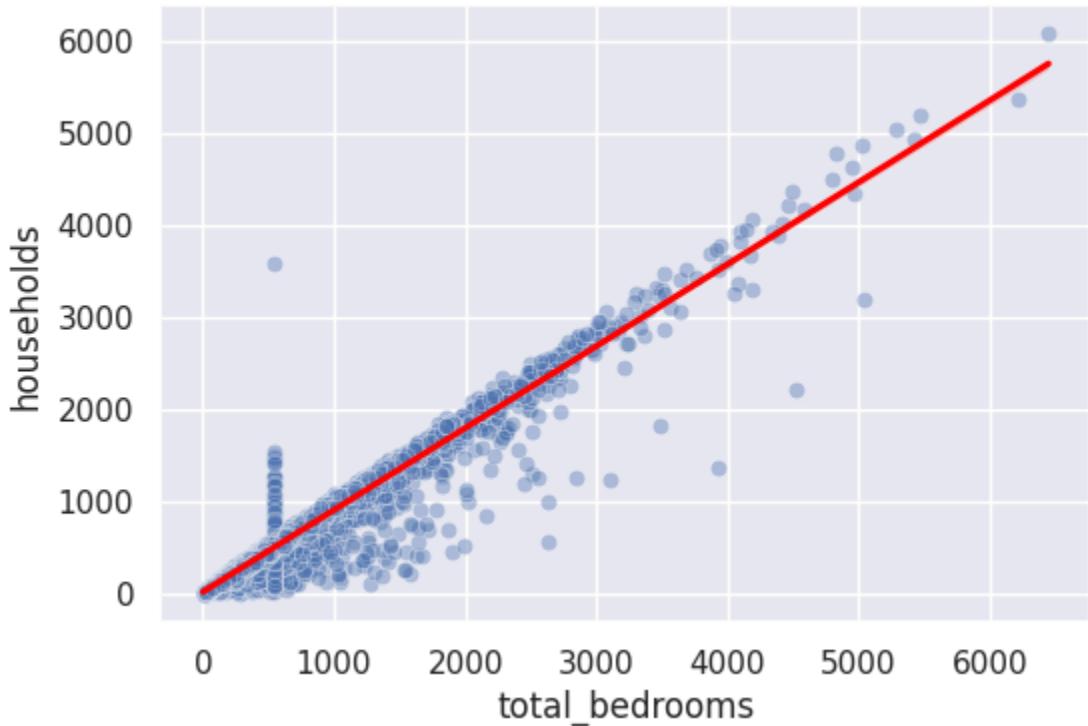
# get top 5 correlated pairs
pairs = (
    corr.where(np.triu(np.ones(corr.shape), k=1).astype(bool))
    .stack()
    .sort_values(ascending=False)
    .head(5)
)

print(pairs)

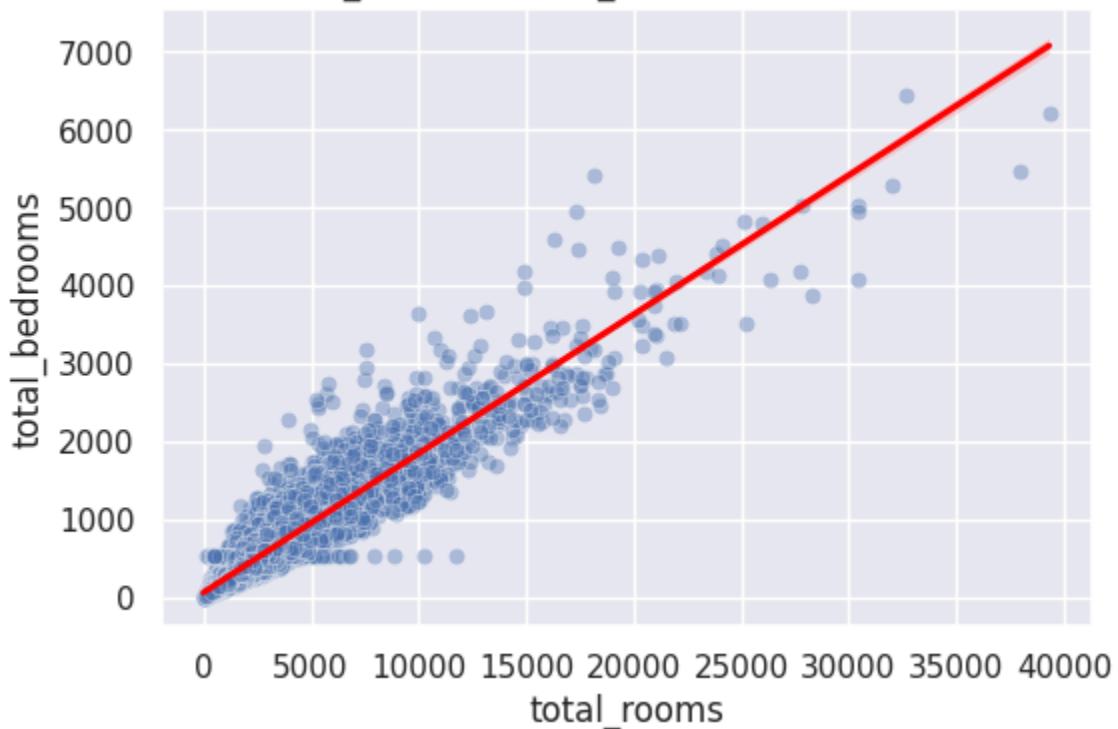
for (col1, col2), val in pairs.items():
    plt.figure(figsize=(6,4))
    sns.scatterplot(x=california_df[col1], y=california_df[col2], alpha=0.4)
    sns.regplot(x=california_df[col1], y=california_df[col2], scatter=False, c
```

```
total_bedrooms    households      0.974725
total_rooms       total_bedrooms  0.927253
longitude         latitude        0.924664
total_rooms       households      0.918484
population        households      0.907222
dtype: float64
```

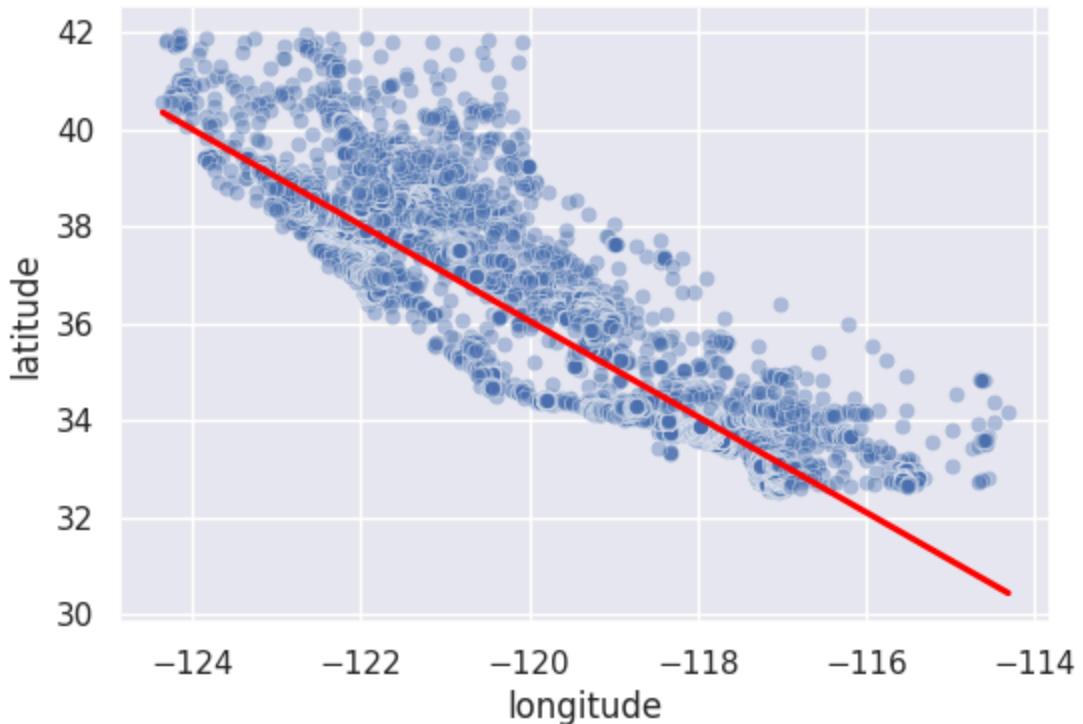
total_bedrooms vs households (corr=0.97)



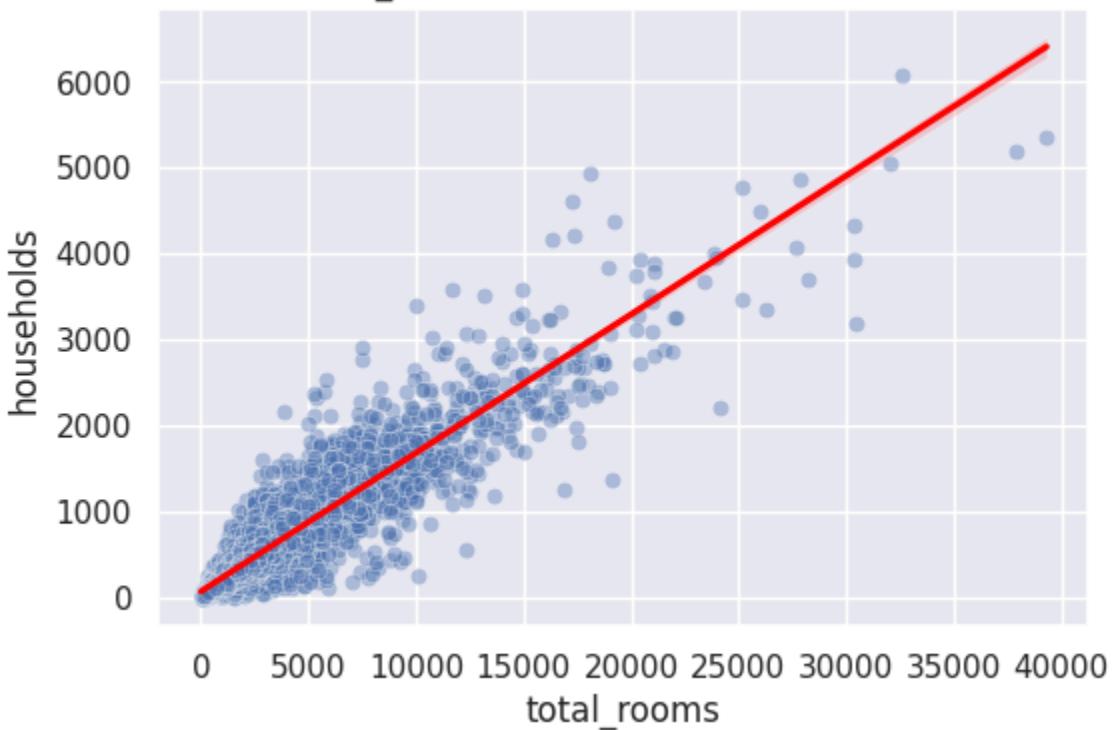
total_rooms vs total_bedrooms (corr=0.93)

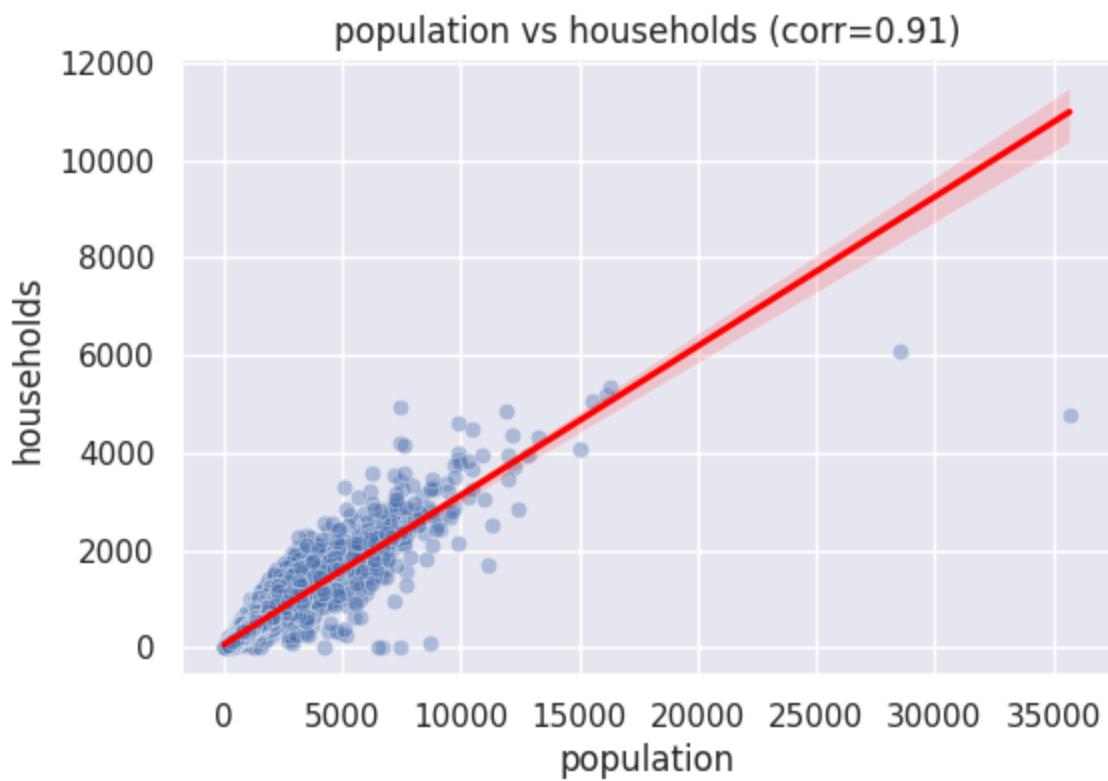


longitude vs latitude (corr=0.92)



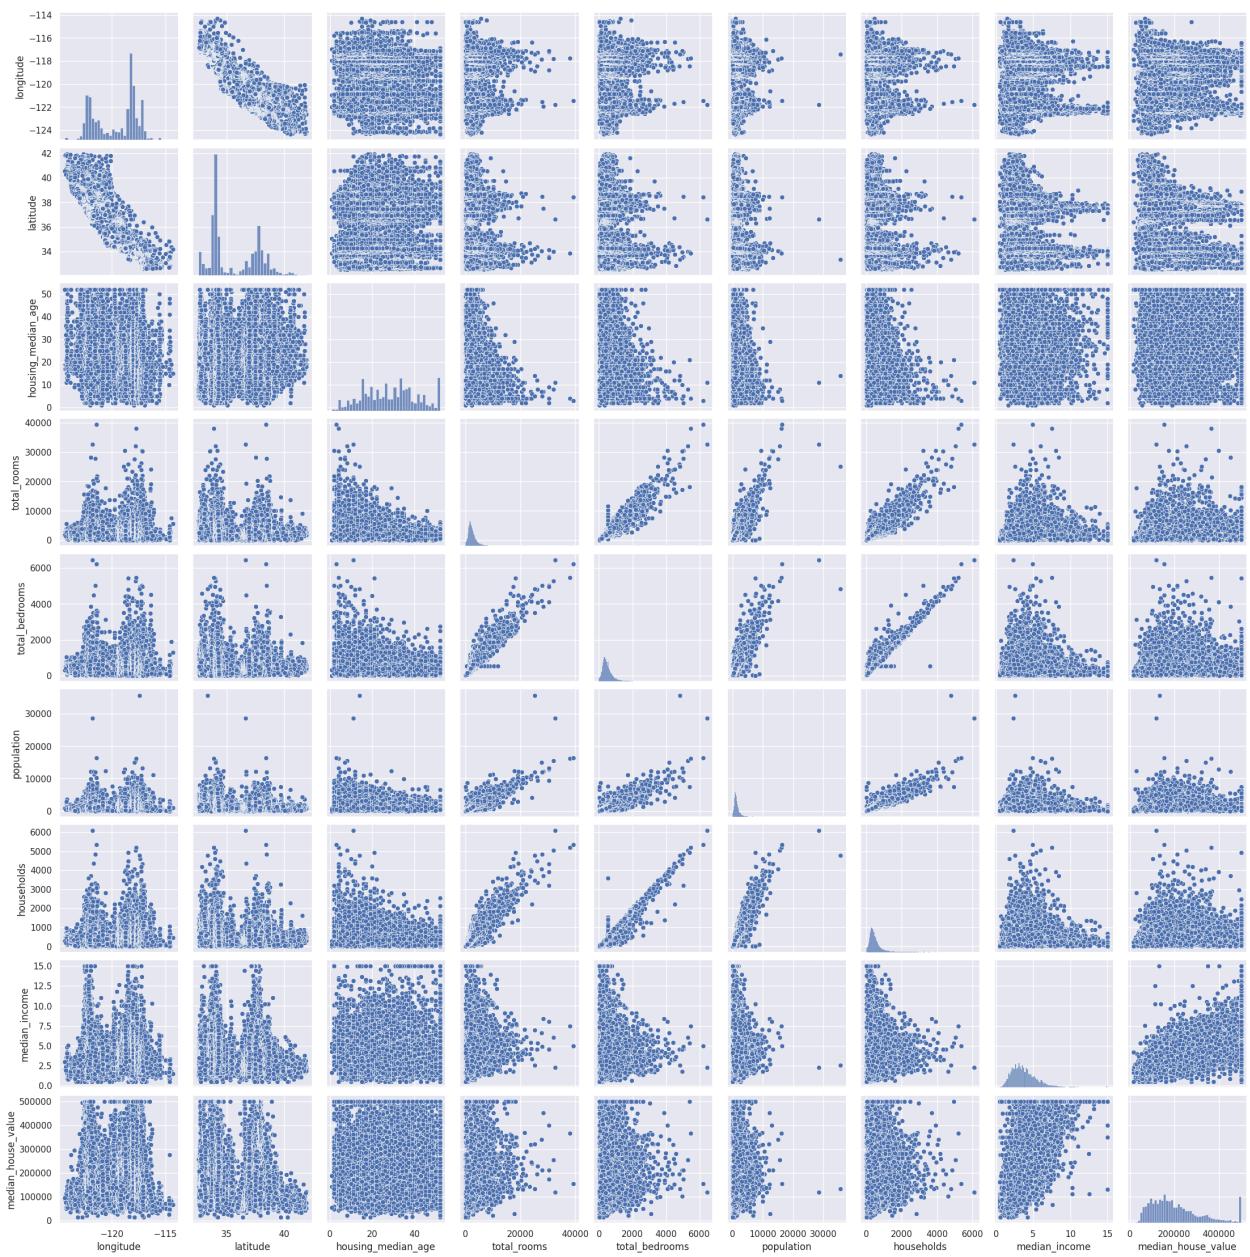
total_rooms vs households (corr=0.92)





```
In [151]: # Get pair values  
plt.figure(figsize=(25,15))  
  
# Plot  
sns.pairplot(california_df)  
  
# Sho  
plt.show()
```

<Figure size 2500x1500 with 0 Axes>



```

In [152]: # Get the longitudinal & latitudinal graph
sns.set(rc={'figure.figsize':(11.7,8.27)})
fig, axes = plt.subplots(1, 2)

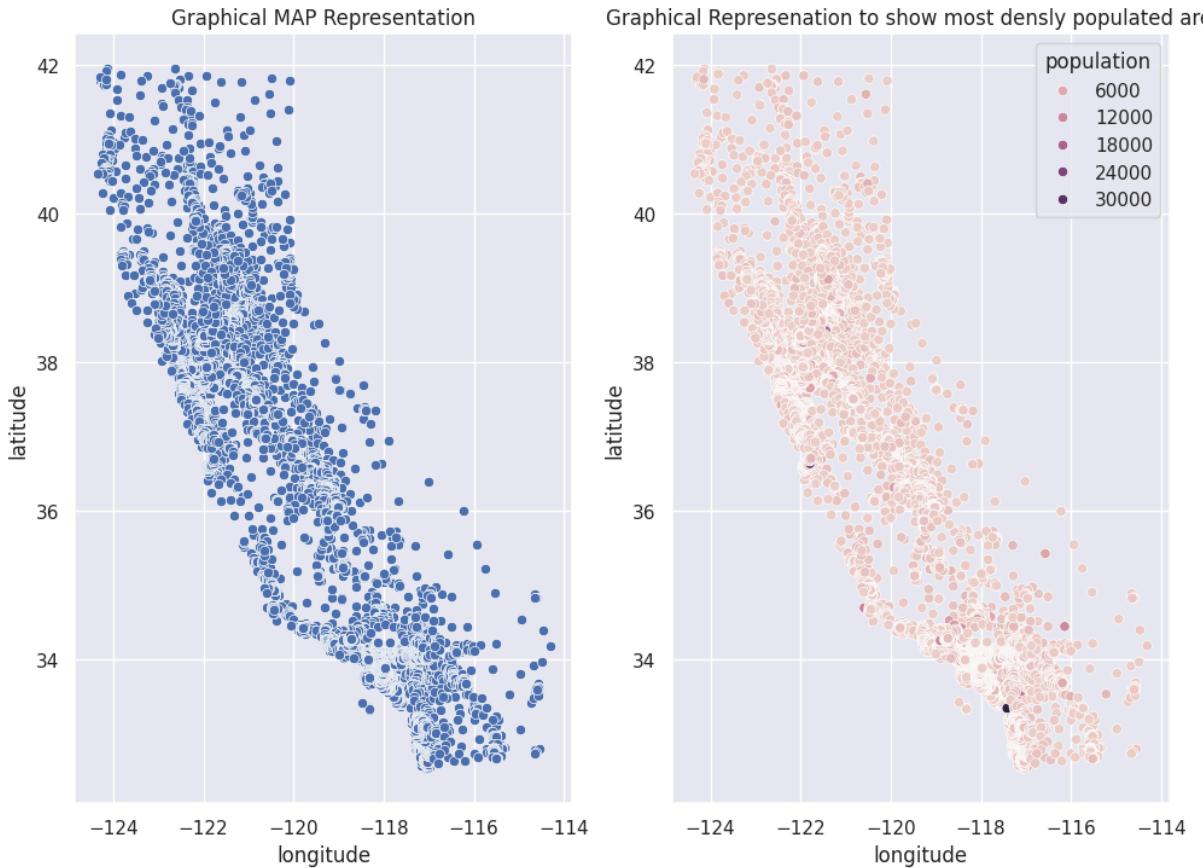
# Plot
sns.scatterplot(data=california_df, x="longitude", y="latitude", ax=axes[0]).set_title("Graphical MAP Representation")
sns.scatterplot(data=california_df, x="longitude", y="latitude", hue="population", ax=axes[1]).set_title("Graphical Representation to show most densely populated areas")

# Show
plt.show()

# ||| ANALYSIS |||

```

```
# We observe that the mostly densely populated areas are near San Diego & Los
```



In [153]:

```
# The relation between the Ocean Proximity and Population - PLOT 1
sns.set(rc={'figure.figsize':(11.7,8.27)})
```

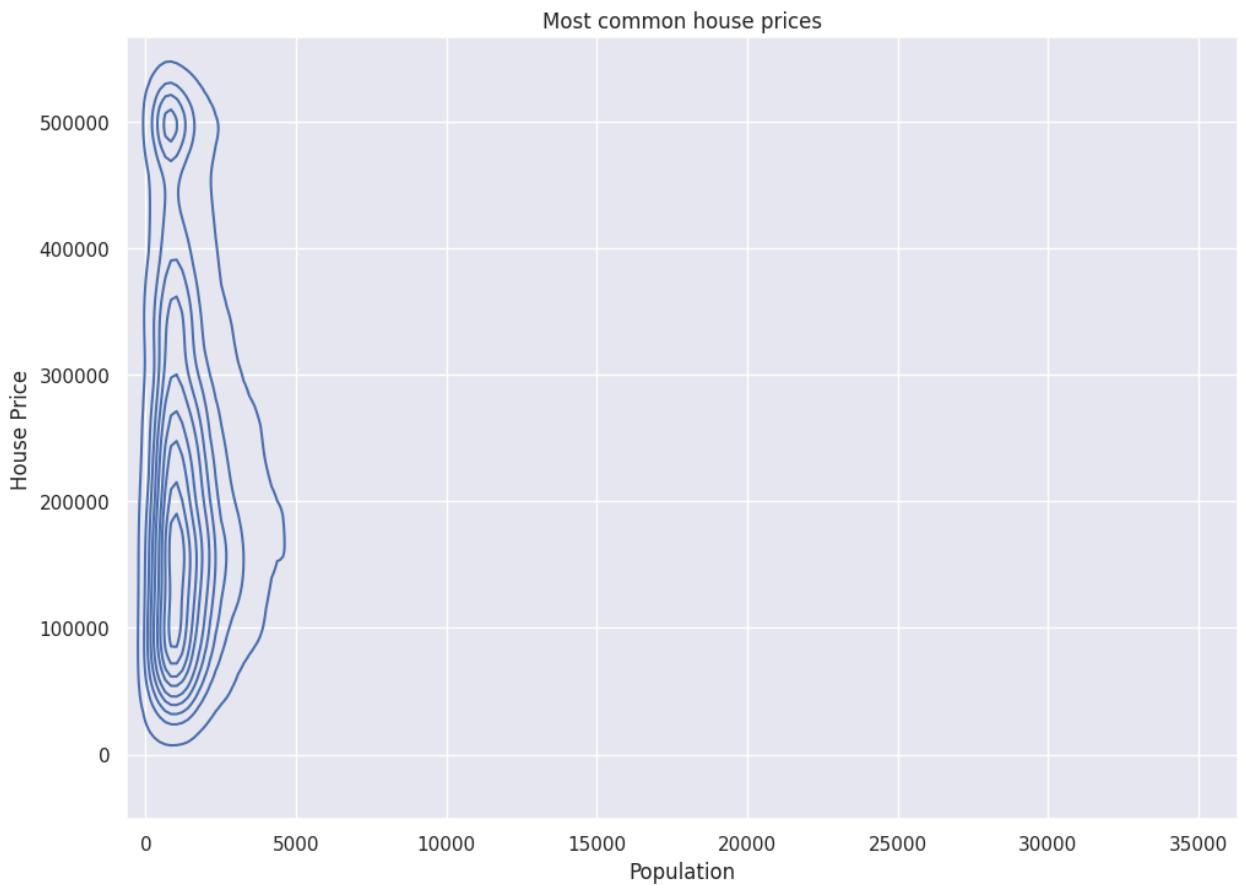
```
# Plot
sns.scatterplot(data=california_df, x="population", y="median_house_value").set(
    title = "How does price of house change with population and ocean proximity",
    xlabel = "Population",
    ylabel = "House Price"
)
# Show
plt.show()
```

```
# |||ANALYSIS|||
# We observe that most of the houses are Inland, with a price between 30,000
```

How does price of house change with population and ocean proxmty - PLOT 1



```
In [154]: # The most common house prices
sns.set(rc={'figure.figsize':(11.7,8.27)})  
  
# Plot
sns.kdeplot(x='population', y='median_house_value', data=california_df).set(
    title = "Most common house prices",
    xlabel = "Population",
    ylabel = "House Price"
)  
  
# Show
plt.show()  
  
# |||ANALYSIS|||
# We observe that mostly house prices in CA, lie between 50,000-280,000.
```



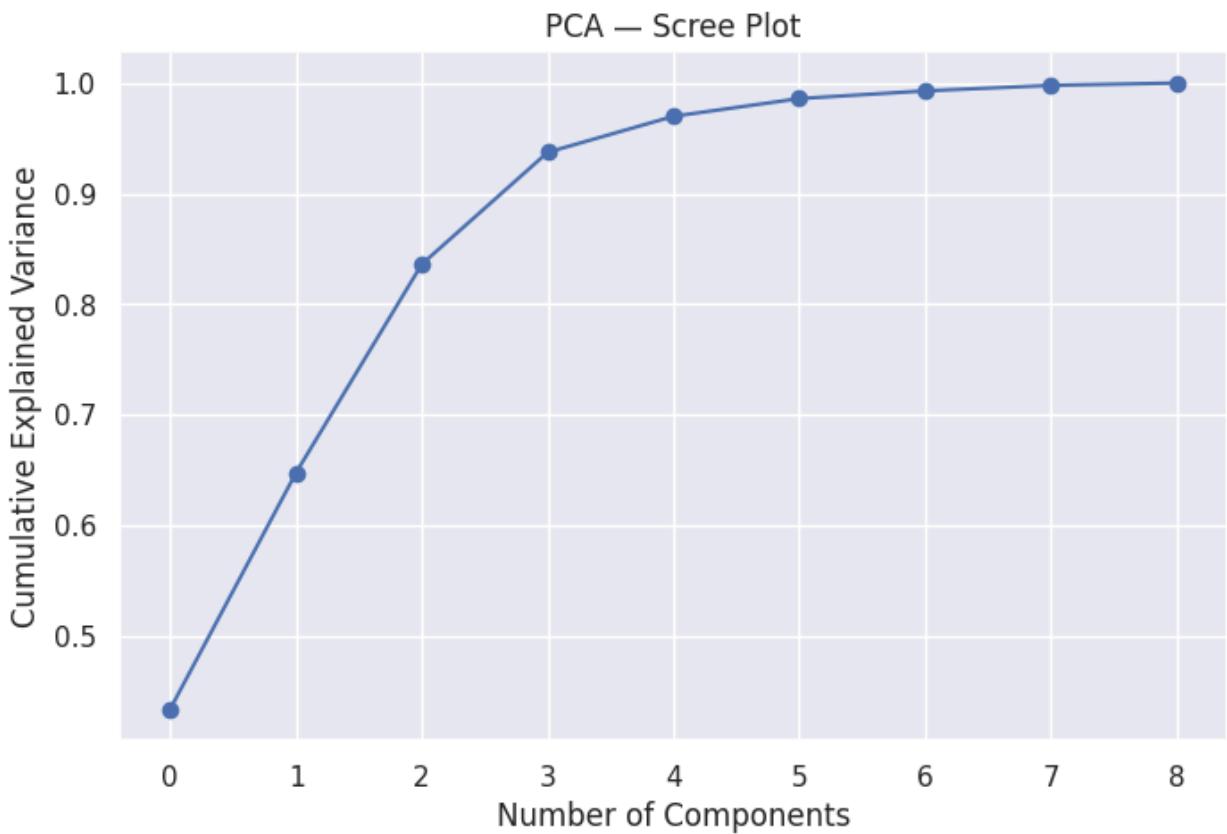
MULTI VARIATE ANALYSIS

```
In [155...]: from sklearn.preprocessing import StandardScaler
          from sklearn.decomposition import PCA

          # Standardize numeric features
          X_scaled = StandardScaler().fit_transform(california_df[numeric_cols].fillna(0))

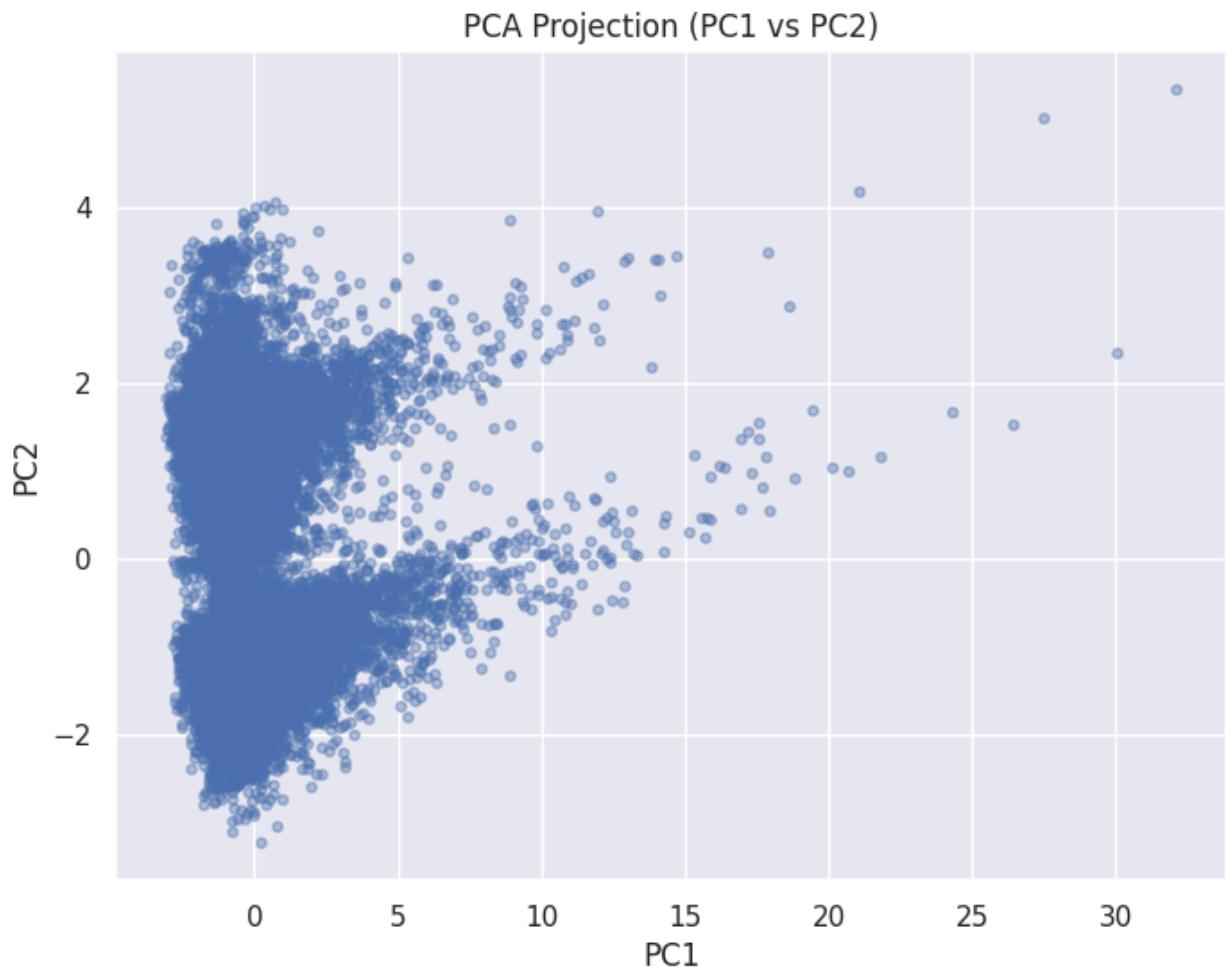
          # PCA
          pca = PCA()
          pca.fit(X_scaled)

          # Explained variance plot (scree)
          plt.figure(figsize=(8,5))
          plt.plot(np.cumsum(pca.explained_variance_ratio_), marker="o")
          plt.xlabel("Number of Components")
          plt.ylabel("Cumulative Explained Variance")
          plt.title("PCA - Scree Plot")
          plt.grid(True)
          plt.show()
```



```
In [156]: X_pca = pca.transform(X_scaled)

plt.figure(figsize=(8,6))
plt.scatter(
    X_pca[:, 0],
    X_pca[:, 1],
    alpha=0.4,
    s=15
)
plt.xlabel("PC1")
plt.ylabel("PC2")
plt.title("PCA Projection (PC1 vs PC2)")
plt.grid(True)
plt.show()
```



SIMPLE LINEAR REGRESSION

```
In [157...]: from sklearn.model_selection import train_test_split
          from sklearn.linear_model import LinearRegression
          from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
In [158...]: df = pd.read_csv("housing[1].csv")
```

```
In [159...]: X = df[["median_income"]]           # independent variable
          y = df["median_house_value"]
```

```
In [160...]: X_train, X_test, y_train, y_test = train_test_split(
              X, y, test_size=0.2, random_state=42
            )
```

```
In [161...]: model = LinearRegression()
          model.fit(X_train, y_train)
```

```
Out[161]: ▾ LinearRegression ⓘ ?  
LinearRegression()
```

```
In [162]: y_pred = model.predict(X_test)
```

```
In [163]: print("Coefficient:", model.coef_[0])  
print("Intercept:", model.intercept_)  
print("MAE:", mean_absolute_error(y_test, y_pred))  
print("R²:", r2_score(y_test, y_pred))
```

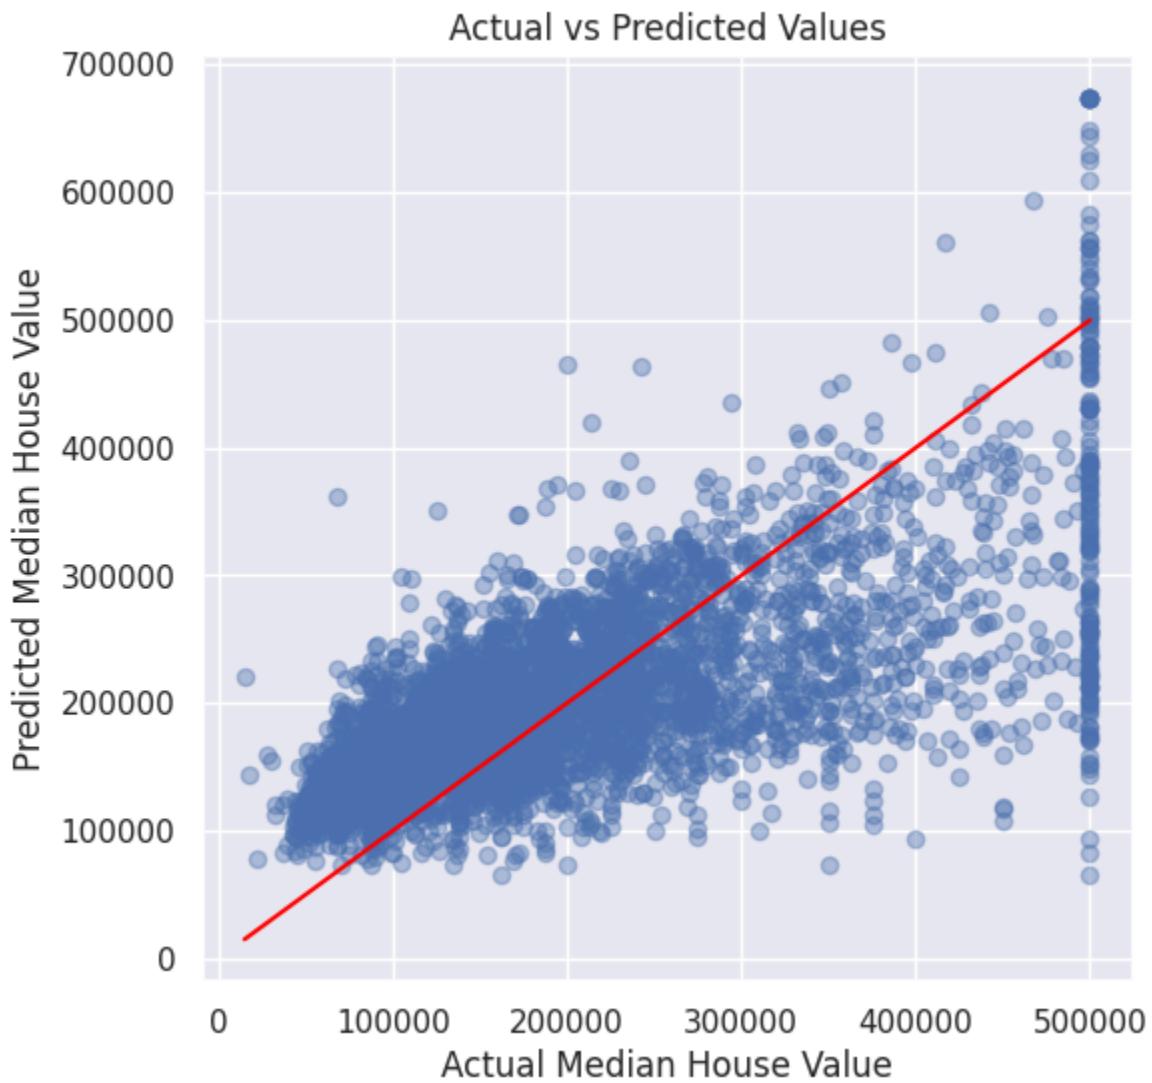
Coefficient: 41933.84939381272
Intercept: 44459.72916907875
MAE: 62990.86530093761
R²: 0.45885918903846656

```
In [164]: plt.figure(figsize=(8,5))  
sns.scatterplot(x=X_test["median_income"], y=y_test, alpha=0.4)  
sns.lineplot(x=X_test["median_income"], y=y_pred, color='red')  
plt.xlabel("Median Income")  
plt.ylabel("Median House Value")  
plt.title("Simple Linear Regression: Median Income vs House Value")  
plt.show()
```



```
In [165]: plt.figure(figsize=(6,6))  
plt.scatter(y_test, y_pred, alpha=0.4)  
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], color="red")
```

```
plt.xlabel("Actual Median House Value")
plt.ylabel("Predicted Median House Value")
plt.title("Actual vs Predicted Values")
plt.show()
```



MULTIPLE LINEAR REGRESSION

```
In [166]: from sklearn.impute import SimpleImputer
```

```
In [167]: numeric_df = df.select_dtypes(include=['number'])
```

```
In [168]: # Set the target variable
target = "median_house_value"           # your dataset's target
```

```
In [169]: # Define X and y
X = numeric_df.drop(columns=[target])    # all numeric features except target
y = numeric_df[target]
```

```
In [170... imputer = SimpleImputer(strategy="median")
X_imputed = imputer.fit_transform(X)
```

```
In [171... X_train, X_test, y_train, y_test = train_test_split(
    X_imputed, y, test_size=0.2, random_state=42
)
```

```
In [172... model = LinearRegression()
model.fit(X_train, y_train)
```

```
Out[172... ▾ LinearRegression ⓘ ?
```

```
LinearRegression()
```

```
In [173... y_pred = model.predict(X_test)
```

```
In [174... print("Intercept:", model.intercept_)
print("\n==== MODEL PERFORMANCE ===")
print("MAE:", mean_absolute_error(y_test, y_pred))
print("R²:", r2_score(y_test, y_pred))
```

```
Intercept: -3578224.234818384
```

```
==== MODEL PERFORMANCE ===
MAE: 51810.483628046866
R²: 0.6138664756435175
```

```
In [175... plt.figure(figsize=(7,6))
sns.scatterplot(x=y_test, y=y_pred, alpha=0.4)
plt.plot([y_test.min(), y_test.max()],
          [y_test.min(), y_test.max()],
          color='red', linewidth=2)

plt.xlabel("Actual Median House Value")
plt.ylabel("Predicted Median House Value")
plt.title("Multiple Linear Regression: Actual vs Predicted")
plt.show()
```

Multiple Linear Regression: Actual vs Predicted

